

Exp. No : 5**HIVE Installation**

1. Download and Extract apache hive 3.1.2

```
lksh@fedora:~$ cd hive
lksh@fedora:~/hive$ ls
bin                examples  lib          RELEASE_NOTES.txt
binary-package-licenses hcatalog LICENSE      scripts
conf              jdbc      NOTICE
```

2. Update HIVE Configurations in .bashrc

```
export HADOOP_CONF_DIR=$HADOOP_HOME/etc/hadoop
export HADOOP_COMMON_LIB_NATIVE_DIR=$HADOOP_HOME/lib/native
export HADOOP_STREAMING=$HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming-3.4.0.jar
export HADOOP_LOG_DIR=$HADOOP_HOME/logs
export PDSH_RCMD_TYPE=ssh

#PIG settings
export PIG_HOME=/home/harithaah/pig
export PATH=$PATH:$PIG_HOME/bin
export PIG_CLASSPATH=$PIG_HOME/conf:$HADOOP_INSTALL/etc/hadoop/
export PIG_CONF_DIR=$PIG_HOME/conf
export JAVA_HOME=/usr/lib/jvm/jdk1.8.0_202
export PATH=$PATH:$JAVA_HOME/bin
export PIG_CLASSPATH=$PIG_CONF_DIR:$PATH
#PIG setting ends

export HIVE_HOME=/home/harithaah/hive
export HIVE_CONF_DIR=$HIVE_HOME/conf
export PATH=$PATH:$HIVE_HOME/bin
export HADOOP_USER_CLASSPATH_FIRST=true
export CLASSPATH=$CLASSPATH:$HIVE_CONF_DIR:$HIVE_HOME/lib/*
```

3. Change directory to apache-hive-3.1.2-bin/conf

```
lksh@fedora:~/hive$ cd conf
lksh@fedora:~/hive/conf$ ls
beeline-log4j2.properties.template  hive-site.xml
hive-default.xml.template           ivysettings.xml
hive-env.sh                         llap-cli-log4j2.properties.template
hive-env.sh.template               llap-daemon-log4j2.properties.template
hive-exec-log4j2.properties.template  parquet-logging.properties
hive-log4j2.properties.template
```

4. Create hive-env.sh

```

GNU nano 7.2                                                                    hive-env.sh
# if [ "$SERVICE" = "cli" ]; then
#   if [ -z "$DEBUG" ]; then
#     export HADOOP_OPTS="$HADOOP_OPTS -XX:NewRatio=12 -Xms10m -XX:MaxHeapFreeRatio=40 -XX:MinHeapFreeRatio=15 -XX:+UseParNewGC -XX:-UseGCOverheadLimit"
#   else
#     export HADOOP_OPTS="$HADOOP_OPTS -XX:NewRatio=12 -Xms10m -XX:MaxHeapFreeRatio=40 -XX:MinHeapFreeRatio=15 -XX:-UseGCOverheadLimit"
#   fi
# fi

# The heap size of the jvm started by hive shell script can be controlled via:
#
# export HADOOP_HEAPSIZE=1024
#
# Larger heap size may be required when running queries over large number of files or partitions.
# By default hive shell scripts use a heap size of 256 (MB). Larger heap size would also be
# appropriate for hive server.

# Set HADOOP_HOME to point to a specific hadoop install directory
# HADOOP_HOME=${bin}/../../hadoop

# Hive Configuration Directory can be controlled by:
# export HIVE_CONF_DIR=

# Folder containing extra libraries required for hive compilation/execution can be controlled by:
# export HIVE_AUX_JARS_PATH=

export HADOOP_HOME=$HADOOP_HOME
export HIVE_CONF_DIR=$HIVE_CONF_DIR

```

5. Install and Change mysql root password

```

Enter password:
Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 8
Server version: 8.0.39 Source distribution

Copyright (c) 2000, 2024, Oracle and/or its affiliates.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql> SELECT user, host, plugin FROM mysql.user WHERE user = 'root';
+-----+-----+-----+
| user | host      | plugin                |
+-----+-----+-----+
| root | localhost | mysql_native_password |
+-----+-----+-----+
1 row in set (0.00 sec)

mysql> ALTER USER 'root'@'localhost' IDENTIFIED WITH 'mysql_native_password' BY 'hayag
reevan';
Query OK, 0 rows affected (0.62 sec)

mysql> FLUSH PRIVILEGES;
Query OK, 0 rows affected (0.12 sec)

mysql>

```

5. Create and configure hive-site.xml

```
GNU nano 7.2      hive-site.xml

<property>
<name>javax.jdo.option.ConnectionURL</name>
<value>jdbc:mysql://localhost/metastore?createDatabaseIfNotExist=true</value>
</property>
<property>
<name>javax.jdo.option.ConnectionDriverName</name>
<value>com.mysql.cj.jdbc.Driver</value>
</property>
<property>
<name>javax.jdo.option.ConnectionUserName</name>
<value>root</value>
</property>

<property>
<name>javax.jdo.option.ConnectionPassword</name>
<value>hayagreevan</value>
</property>
<property>
<name>datanucleus.autoCreateSchema</name>
<value>true</value>
</property>
<property>
<name>datanucleus.fixedDatastore</name>
<value>true</value>
</property>
<property>
<name>datanucleus.autoCreateTables</name>
<value>True</value>
</property>
</configuration>

^G Help      ^O Write Out  ^W Where Is   ^K Cut        ^T Execute    ^C Location
^X Exit      ^R Read File  ^\ Replace    ^U Paste      ^J Justify    ^_ Go To Line
```

6. Download and Move mysql java connector to apache-hive-3.1.2-bin/lib

```

lksh@fedora: ~/hive/conf$ nano hive-env.sh
lksh@fedora: ~/hive/conf$ cd ..
lksh@fedora: ~/hive$ cd lib
lksh@fedora: ~/hive/lib$ ls
accumulo-core-1.7.3.jar      datanucleus-api-jdo-4.2.4.jar      hive-hcatalog-server-extensions-3.1.2.jar  javax.jdo-3.2.0-m3.jar
accumulo-fate-1.7.3.jar      datanucleus-core-4.1.17.jar         hive-hplsql-3.1.2.jar                    javax.servlet-api-3.1.0.jar
accumulo-start-1.7.3.jar     datanucleus-rdbms-4.1.19.jar        hive-jdbc-3.1.2.jar                     javax.servlet.jsp-2.3.2.jar
accumulo-trace-1.7.3.jar     derby-10.14.1.0.jar                hive-jdbc-handler-3.1.2.jar              javax.ws.rs-api-2.0.1.jar
aircompressor-0.10.jar      disruptor-3.3.6.jar                hive-kryo-registrator-3.1.2.jar          javolution-5.5.1.jar
ant-1.9.1.jar               druid-hdfs-storage-0.12.0.jar       hive-llap-client-3.1.2.jar               jaxb-api-2.2.11.jar
ant-launcher-1.9.1.jar      esri-geometry-api-2.0.0.jar         hive-llap-common-3.1.2.jar               jcodings-1.0.18.jar
antlr-runtime-4.5.jar       flatbuffers-1.2.0-3f79e855.jar     hive-llap-common-3.1.2-tests.jar         jcommander-1.32.jar
aopalliance-repackaged-2.5.0-b32.jar  findbugs-annotations-1.3.9-1.jar  hive-llap-ext-client-3.1.2.jar           jdo-api-3.0.1.jar
apache-curator-2.12.0.pom   gson-2.2.4.jar                    hive-llap-tez-3.1.2.jar                  jersey-client-2.25.1.jar
apache-jsp-9.3.20.v20170531.jar  groovy-all-2.4.11.jar             hive-metastore-3.1.2.jar                 jersey-common-2.25.1.jar
apache-jstl-9.3.20.v20170531.jar  guava-19.0.jar                    hive-serde-3.1.2.jar                    jersey-container-servlet-core-2.25.1.jar
arrow-format-0.8.0.jar      hbase-client-2.0.0-alpha4.jar       hive-service-3.1.2.jar                   jersey-guava-2.25.1.jar
arrow-memory-0.8.0.jar      hbase-common-2.0.0-alpha4.jar       hive-shims-0.23-3.1.2.jar                jersey-media-jaxb-2.25.1.jar
arrow-vector-0.8.0.jar      hbase-common-2.0.0-alpha4-tests.jar  hive-shims-0.23-3.1.2.jar                jersey-server-2.25.1.jar
asm-5.0.1.jar               hbase-hadoop2-compat-2.0.0-alpha4.jar  hive-shims-3.1.2.jar                     jettison-1.1.jar
asm-commons-5.0.1.jar       hbase-hadoop2-compat-2.0.0-alpha4-tests.jar  hive-shims-common-3.1.2.jar              jetty-annotations-9.3.20.v20170531.jar
asm-tree-5.0.1.jar          hbase-hadoop-compat-2.0.0-alpha4.jar  hive-shims-scheduler-3.1.2.jar           jetty-client-9.3.20.v20170531.jar
audience-annotations-0.5.0.jar  hbase-http-2.0.0-alpha4.jar         hive-standalone-metastore-3.1.2.jar      jetty-http-9.3.20.v20170531.jar
avatica-1.11.0.jar          hbase-mapreduce-2.0.0-alpha4.jar     hive-storage-api-2.7.0.jar               jetty-io-9.3.20.v20170531.jar
avro-1.7.7.jar              hbase-mapreduce-2.0.0-alpha4.jar     hive-streaming-3.1.2.jar                 jetty-jas-9.3.20.v20170531.jar
bonecp-0.8.0.RELEASE.jar     hbase-metrics-api-2.0.0-alpha4.jar  hive-testutils-3.1.2.jar                 jetty-jas-9.3.20.v20170531.jar
calcite-core-1.16.0.jar      hbase-metrics-2.0.0-alpha4.jar       hive-upgrade-acid-3.1.2.jar              jetty-jndi-9.3.20.v20170531.jar
calcite-druid-1.16.0.jar     hbase-prefix-tree-2.0.0-alpha4.jar  hive-vector-code-gen-3.1.2.jar            jetty-plus-9.3.20.v20170531.jar
calcite-linq4j-1.16.0.jar    hbase-procedure-2.0.0-alpha4.jar     hk2-locator-2.5.0-b32.jar                jetty-rewrite-9.3.20.v20170531.jar
commons-cli-1.2.jar          hbase-protocol-2.0.0-alpha4.jar      hk2-utils-2.5.0-b32.jar                  jetty-runner-9.3.20.v20170531.jar
commons-codec-1.7.jar        hbase-protocol-shaded-2.0.0-alpha4.jar  hppc-0.7.2.jar                           jetty-schemas-3.1.jar
commons-collections4-4.1.jar  hbase-replication-2.0.0-alpha4.jar  htrace-core-3.2.0-incubating.jar          jetty-security-9.3.20.v20170531.jar
commons-compiler-2.7.6.jar    hbase-server-2.0.0-alpha4.jar        httpclient-4.5.2.jar                     jetty-server-9.3.20.v20170531.jar
commons-compress-1.9.jar     hbase-shaded-netty-1.0.1.jar         httpcore-4.4.4.jar                       jetty-servlet-9.3.20.v20170531.jar
commons-crypto-1.0.0.jar     hbase-shaded-protobuf-1.0.1.jar      ivy-2.4.0.jar                             jetty-util-9.3.20.v20170531.jar
commons-dbc-1.4.jar          HikariCP-2.6.1.jar                  jackson-annotations-2.9.5.jar             jetty-webapp-9.3.20.v20170531.jar
commons-io-2.4.jar           hive-accumulo-handler-3.1.2.jar       jackson-core-2.9.5.jar                   jetty-xml-9.3.20.v20170531.jar
commons-lang-2.6.jar          hive-beeline-3.1.2.jar                jackson-core-asl-1.9.13.jar               jline-2.12.jar
commons-lang-3.2.jar          hive-classification-3.1.2.jar         jackson-databind-2.9.5.jar                joda-time-2.9.9.jar
commons-logging-1.0.4.jar     hive-cli-3.1.2.jar                   jackson-dataformat-smile-2.9.5.jar         joni-2.1.11.jar
commons-math-2.1.jar          hive-common-3.1.2.jar                 jackson-dataformat-smile-2.9.5.jar         jpam-1.1.jar
commons-math3-3.6.1.jar       hive-contrib-3.1.2.jar                jackson-mapper-asl-1.9.13.jar             json-1.8.jar
commons-pool-1.5.4.jar        hive-druid-handler-3.1.2.jar          jamon-runtime-2.3.1.jar                   jsr305-3.0.0.jar
commons-vfs2-2.1.jar          hive-exec-3.1.2.jar                  janino-2.7.6.jar                          jta-1.1.jar
curator-client-2.12.0.jar     hive-hbase-handler-3.1.2.jar          javassist-3.20.0-GA.jar                   libfb303-0.9.3.jar
curator-framework-2.12.0.jar  hive-hcatalog-core-3.1.2.jar          javax.annotation-api-1.2.jar               libthrift-0.9.3.jar
curator-recipes-2.12.0.jar    hive-hcatalog-server-3.1.2.jar        javax.inject-2.5.0-b32.jar                log4j-1.2-api-2.10.0.jar
lksh@fedora: ~/hive/lib$

```

7. Execute schematool -initSchema -dbType mysql

```

lf4j-reload4j-1.7.36.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Metastore connection URL:      jdbc:mysql://localhost/metastore?createDatabaseIfNo
tExist=true
Metastore Connection Driver :   com.mysql.cj.jdbc.Driver
Metastore connection User:     root
Starting metastore schema initialization to 3.1.0
Initialization script hive-schema-3.1.0.mysql.sql

```

8. Start hive

```
curator-recipes-2.12.0.jar      hive-hcatalog-core-3.1.2.jar      javax.inject-2.5.0-b32.jar      log4j
lksh@fedora:~/hive/lib$ hive
which: no hbase in (/home/lksh/.local/bin:/home/lksh/bin:/usr/local/bin:/usr/local/sbin:/usr/bin:/usr/sbin:/usr/lib/jvm/jdk-1.8-oracle-x64/bin:/usr/sbin:/usr/bin:/usr/sbin:/usr/lib/jvm/jdk-1.8-oracle-x64/bin:/home/lksh/hadoop/bin:/home/lksh/hadoop/sbin:/home/lksh/pig/bin:/home/lksh/hive/bin)
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/lksh/hive/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/lksh/hadoop/share/hadoop/common/lib/slf4j-reload4j-1.7.36.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Hive Session ID = 56312152-5bcf-4fa4-9f9e-b68d787b2898

Logging initialized using configuration in jar:file:/home/lksh/hive/lib/hive-common-3.1.2.jar!/hive-log4j2.properties Async: true
Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez, or mr).
hive>
```